

# Spark partitions

1. what happens if there are less partitions than executors?

1 point

*Mark only one oval.*

- ☐ Spark will not be able to use all the executors to run in parallel, and the computation result will be correct
- ☐ Spark will not be able to use all the executors to run in parallel, and the computation result will be wrong
- ☐ it has no effect

2. what happens if there are more partitions than executors?

1 point

*Mark only one oval.*

- ☐ Spark will not be able to use all the executors to run in parallel
- ☐ This situation is impossible
- ☐ All executors will work, and some of them will handle more than one partition

3. Who is responsible to partition the data?

1 point

*Mark only one oval.*

- ☐ First executor
- ☐ the Driver node
- ☐ executors voting
- ☐ The cluster manager

4. what is the criterion to decide if a row is in a certain partition?

1 point

*Mark only one oval.*

- ☐ assigning a row to a partition is random
- ☐ an object hash is computed for the row, and the value is mapped to one of the partitions
- ☐ first K rows go to the first partition, the next K to the second etc.

5. can an RDD be partitioned by columns only (vertical partitioning)?

1 point

*Mark only one oval.*

- ☐ yes
- ☐ no
- ☐ cannot decide based on what we learned so far

6. Is partitioning automatic? Can I influence it?

1 point

*Mark only one oval.*

- ☐ yes, there is no way to influence it
- ☐ yes, we can influence mapping rows to partitions
- ☐ no. The user has to supply a function that will tell Spark how to map each row to a partition

---

This content is neither created nor endorsed by Google.

Google Forms

